

JING SU

DATA & AI SYSTEMS | CLOUD & DATA ARCHITECTURE | GENAI / MLOPS
Champaign, IL | (217) 281-2782 | sujingmaster@gmail.com | linkedin.com/in/jingsu-ds

PROFESSIONAL SUMMARY

Data & AI systems professional with 5+ years of experience building production data platforms, AI-enabled automation, and cloud-based data and ML systems. Brings a strong foundation in Python/SQL, AWS data engineering, cross-system integration, and GenAI/ML, with increasing ownership of solution design, technical delivery, and business-to-technical translation. Built a production LLM-powered web-intelligence pipeline processing 25K+ supplier domains and extracting 375K+ products and 222K+ deduplicated ingredients for roughly \$90 in AI cost; also automated a lead-qualification workflow from a 5–7-person, six-week process to a single-machine run completed in about five days. Partners across business, engineering, and leadership stakeholders and mentors junior engineers while driving ambiguous problems from requirements through implementation and validation. Holds an MCS from UIUC with a 4.0 GPA.

CORE TECHNICAL SKILLS

Data Engineering	Python, SQL, PySpark, Spark SQL; ETL/ELT, data pipelines, data modeling, data quality and reconciliation, API integration, batch & streaming processing
AI / ML / NLP / GenAI	Machine Learning, Natural Language Processing (NLP), forecasting, model evaluation, scikit-learn, TensorFlow; OpenAI API, embeddings, semantic search / vector retrieval, RAG, batch inference pipelines, MLOps
Cloud & Platform	AWS (S3, EC2, Lambda, Glue, Athena, Redshift, DynamoDB, RDS, SageMaker, Step Functions, EKS, ECR, CloudWatch); GCP (BigQuery, Pub/Sub, Dataflow, Cloud Functions, Gemini); Docker, Kubernetes, Git/GitHub
Data & Enterprise Systems	MongoDB, PostgreSQL, MySQL, SQL Server, Redis; NetSuite, HubSpot, Magento, Akeneo; Crawl4AI, Playwright, Scrapy, Tableau

PROFESSIONAL EXPERIENCE

Green Wave Ingredients (Ingredients Online)	Orange County, CA Remote
Data Engineer	May 2024 - Present
• Built a production web-intelligence and ingredient-extraction pipeline on Crawl4AI, the OpenAI Batch API, cost-optimized EC2, and S3, processing 25,733 supplier domains into 375,733 products and 222,035 deduplicated ingredients for roughly \$90 in LLM processing cost; S3-backed checkpointing and stale-batch recovery prevented duplicate processing and spend. • Replaced a multi-person manual company and lead qualification process with a single-machine automated workflow that completes in about five days, cutting execution time and coordination overhead. • Design and operate a cross-system data platform spanning NetSuite, HubSpot, Magento, Akeneo, MongoDB, GA4/BigQuery, and AWS (S3, EC2, Redshift, Athena) that powers operational analytics, product intelligence, and AI-enabled workflows. • Own technical delivery for Flowmerce AI, an AI-enabled inventory planning and demand forecasting initiative; delivered the trusted-data foundation for a 30-SKU MVP with reproducible inventory reconciliation and tiered data-quality eligibility controls, and launched the demand-forecasting phase. • Implemented OpenAI- and embedding-based product classification and retrieval workflows, improving categorization accuracy by 40% and cutting query time by 35% through optimized vectorization and storage. • Mentor 2-4 interns and junior engineers annually across onboarding, task assignment, code review, technical training, and project planning, while partnering across business, engineering, and leadership stakeholders.	
Full Stack Data Scientist	May 2021 - May 2024
• Designed a scalable AWS data warehouse on S3, DynamoDB, Glue, and Athena, improving query performance by 30% while reducing associated cloud cost by 40%. • Optimized AWS Glue and PySpark ETL pipelines for billion-row workloads, improving runtime by 80% and reducing processing cost by 90% across targeted jobs. • Developed automated data-error detection and correction controls for the data lake, reducing data errors by 98% and improving the reliability of downstream analytics and operational decisions. • Integrated NetSuite SuiteTalk, Google Analytics, BigQuery, Search Console, and REST APIs into automated data workflows, cutting BI reporting turnaround by 50%; built recommendation and demand-forecasting solutions with AWS Personalize, SageMaker, and scikit-learn.	
Data Analyst	Nov 2020 - May 2021
• Built Tableau dashboards and Python-based validation and reporting workflows for sales, product, and marketing data, improving transparency and reducing manual quality checks. • Implemented Google Analytics tracking and translated cross-functional business questions into recurring analyses and reporting.	

EDUCATION & CERTIFICATION

University of Illinois Urbana-Champaign

Urbana-Champaign, IL | May 2026

Master of Computer Science (MCS) | GPA: 4.0/4.0

Selected coursework: Cloud Computing Applications, Applied Machine Learning, Natural Language Processing, Database Systems, Data Cleaning & Data Curation

University of Illinois Urbana-Champaign

Urbana-Champaign, IL | 2020

Master of Science in Economics - Specialization: Advanced Econometrics

Jilin University

China | 2016

Dual Bachelor's Degrees in Finance and International Politics

Certification: Oracle Cloud Infrastructure 2025 Generative AI Professional

SELECTED HANDS-ON AI & CLOUD ARCHITECTURE PROJECTS

Graduate engineering projects demonstrating architecture-level capabilities; separate from professional production experience.

AWS End-to-End AI Model Lifecycle Automation Platform

- Built an automated ML lifecycle on SageMaker, Step Functions, Lambda, S3, and CloudWatch spanning preprocessing, training, evaluation, conditional model promotion, and production endpoint deployment.
- Implemented drift-triggered retraining and champion/challenger evaluation with promotion gates, monitoring, and workflow orchestration for a governed model lifecycle.

Cloud-Native AI Inference Platform on Amazon EKS

- Designed and deployed a containerized AI serving architecture using EKS, Docker/ECR, Kubernetes Jobs, Envoy, and Flask to separate routing, API, and asynchronous inference layers.
- Applied namespaces, RBAC/ServiceAccounts, ResourceQuota, LoadBalancer services, and production-style cloud networking controls to isolate workloads and support governed execution.

Real-Time Enterprise RAG Customer Support Architecture on GCP

- Designed a real-time RAG architecture connecting enterprise knowledge, embeddings and vector retrieval, Gemini, BigQuery, Cloud Functions, Pub/Sub, and Dataflow for grounded customer-support intelligence.
- Integrated vector retrieval and grounded generation with event-driven streaming workflows to connect AI responses with real-time business data and operational analytics.